Snowflake Certification

Complete MCQ Compilation with Correct Answers

Compiled from all 5 uploaded PDFs — 248 unique questions total

Part 1: SnowPro Core Mock Exam (99 Qs) | Part 2: Scenario-Based Practice (60 Qs) | Part 3: Concept Flashcard MCQs (89 Qs)

Contents

● Part 1 — SnowPro Core: 100 Mock Exam Questions (source: SNOWPRO_1.pdf) — 99 questions captured

● Part 2 — Snowflake Scenario-Based Practice Questions (source: snowflake_test.pdf / certification_real.pdf) — 60 questions

● Part 3 — Snowflake Concept & Exam-Angle Flashcard MCQs (source: Snowflake Certification.pdf) — 89 unique questions

Note: *The source "Snowflake.pdf" (detailed notes deck) contains no MCQs — it is reference/concept material only, so no questions were extracted from it. The SnowPro Core file is titled "100 Mock Exam Questions" but the source PDF itself only contains Questions 1–99 (Question 100 is not present in the file).*

Part 1 — SnowPro Core: 100 Mock Exam Questions

Source file: SNOWPRO_1.pdf

SECTION 1: Snowflake Architecture & Concepts (1–15)

Q1. What is Snowflake's architecture called?

A. Shared disk
B. Shared memory
C. Shared nothing ✔
D. Hybrid shared disk
Correct Answer: C. Shared nothing

Q2. Which Snowflake layer stores metadata?

A. Compute
B. Storage
C. Cloud services ✔
D. Virtual warehouse
Correct Answer: C. Cloud services

Q3. Snowflake automatically manages which of the following?

A. Indexes
B. Partitioning
C. Vacuum
D. All of the above ✔
Correct Answer: D. All of the above

Q4. What unit is Snowflake storage billed in?

A. Credits
B. Compute hours
C. TB/month ✔
D. Requests
Correct Answer: C. TB/month

Q5. What object executes SQL queries?

A. Database
B. Schema
C. Table
D. Virtual warehouse ✔

Correct Answer: D. Virtual warehouse

Q6. What happens when a warehouse is suspended?

A. Queries fail
B. Data is unavailable
C. No credits consumed ✔
D. Storage is paused
Correct Answer: C. No credits consumed

Q7. Which edition supports multi-cluster warehouses?

A. Standard
B. Enterprise ✔
C. Business Critical
D. VPS
Correct Answer: B. Enterprise

Q8. What determines query concurrency?

A. Table size
B. Warehouse size ✔
C. Database size
D. Role hierarchy
Correct Answer: B. Warehouse size

Q9. What happens when storage is scaled?

A. Manual action required
B. Credits are consumed
C. Automatic and seamless ✔
D. Warehouse restarts
Correct Answer: C. Automatic and seamless

Q10. Which service manages authentication?

A. Storage layer
B. Compute layer
C. Cloud services ✔
D. Virtual warehouse
Correct Answer: C. Cloud services

Q11. What is the default warehouse state?

A. Suspended ✔
B. Running
C. Resizing
D. Disabled
Correct Answer: A. Suspended

Q12. Snowflake automatically encrypts data using:

A. AES-128
B. AES-192
C. AES-256 ✔
D. RSA
Correct Answer: C. AES-256

Q13. What enables workload isolation?

A. Databases
B. Schemas
C. Virtual warehouses ✔
D. Roles
Correct Answer: C. Virtual warehouses

Q14. Which component handles query optimization?

A. Storage
B. Cloud services ✔
C. Client
D. Driver
Correct Answer: B. Cloud services

Q15. What enables independent scaling of compute?

A. Separation of storage & compute ✔
B. Clustering keys
C. Result cache
D. Fail-safe
Correct Answer: A. Separation of storage & compute

SECTION 2: Warehouses & Resource Management (16–30)

Q16. What happens when AUTO_RESUME is enabled?

A. Warehouse resumes on schedule

B. Warehouse resumes when query arrives ✔

C. Warehouse resumes on data load

D. Warehouse resumes after resize

Correct Answer: B. Warehouse resumes when query arrives

Q17. What is the minimum billing duration for a warehouse?

A. 1 second

B. 10 seconds

C. 30 seconds

D. 60 seconds ✔

Correct Answer: D. 60 seconds

Q18. What setting controls idle shutdown time?

A. AUTO_RESUME

B. AUTO_SUSPEND ✔

C. MAX_CLUSTER_COUNT

D. SCALING_POLICY

Correct Answer: B. AUTO_SUSPEND

Q19. What does MAX_CLUSTER_COUNT control?

A. Query timeout

B. Maximum concurrent queries

C. Upper limit of clusters ✔

D. Storage usage

Correct Answer: C. Upper limit of clusters

Q20. What happens if all clusters are busy?

A. Queries fail

B. Queries are queued ✔

C. Warehouse restarts

D. Storage scales

Correct Answer: B. Queries are queued

Q21. Which warehouse size uses the fewest credits?

A. X-SMALL ✔

B. SMALL

C. MEDIUM

D. LARGE

Correct Answer: A. X-SMALL

Q22. Scaling Policy "Economy" does what?

A. Adds clusters aggressively

B. Adds clusters conservatively ✔

C. Disables scaling

D. Doubles warehouse size

Correct Answer: B. Adds clusters conservatively

Q23. Which feature limits credit consumption?

A. Budget

B. Resource Monitor ✔

C. Network Policy

D. Warehouse Policy

Correct Answer: B. Resource Monitor

Q24. What happens when a resource monitor limit is exceeded?

A. Warehouse scales up

B. Queries slow down

C. Warehouse is suspended ✔

D. Storage is locked

Correct Answer: C. Warehouse is suspended

Q25. Who can assign a resource monitor?

A. ACCOUNTADMIN ✔

B. SYSADMIN

C. SECURITYADMIN

D. Any role

Correct Answer: A. ACCOUNTADMIN

Q26. What metric shows queued queries?

A. Bytes scanned

B. Compilation time

C. Query load

D. Warehouse load ✔

Correct Answer: D. Warehouse load

Q27. Warehouse resizing affects:

A. Only new queries ✔
B. Only running queries
C. Both
D. Neither
Correct Answer: A. Only new queries

Q28. What controls concurrency scaling?

A. Warehouse size ✔
B. Query complexity
C. Auto-clustering
D. Streams
Correct Answer: A. Warehouse size

Q29. How are credits billed for multi-cluster warehouses?

A. Per warehouse
B. Per cluster ✔
C. Per query
D. Per user
Correct Answer: B. Per cluster

Q30. Which warehouse feature improves burst workloads?

A. Clustering
B. Result cache
C. Multi-cluster ✔
D. Secure views
Correct Answer: C. Multi-cluster

SECTION 3: Data Loading & Unloading (31–50)

Q31. Which command loads data into Snowflake?

A. INSERT
B. COPY INTO ✔
C. LOAD DATA
D. IMPORT
Correct Answer: B. COPY INTO

Q32. Which format is columnar?

A. CSV
B. JSON
C. XML
D. Parquet ✔
Correct Answer: D. Parquet

Q33. What does FORCE=TRUE do?

A. Skips errors
B. Reloads files ✔
C. Deletes data
D. Compresses files
Correct Answer: B. Reloads files

Q34. What stage type is managed by Snowflake?

A. External
B. Internal ✔
C. Named
D. Directory
Correct Answer: B. Internal

Q35. Snowpipe is used for:

A. Batch loads
B. Streaming ingestion ✔
C. Query acceleration
D. Replication
Correct Answer: B. Streaming ingestion

Q36. What triggers Snowpipe?

A. Schedule
B. API call
C. Event notification ✔
D. Warehouse
Correct Answer: C. Event notification

Q37. COPY INTO without FORCE loads:

A. All files
B. Only new files ✔
C. Random files

D. Largest files

Correct Answer: B. Only new files

Q38. Which file format preserves schema?

A. CSV

B. JSON

C. Parquet ✔

D. TXT

Correct Answer: C. Parquet

Q39. Validation_mode = RETURN_ERRORS does what?

A. Loads data

B. Deletes bad rows

C. Returns load errors ✔

D. Stops execution

Correct Answer: C. Returns load errors

Q40. What command unloads data?

A. EXPORT

B. UNLOAD

C. COPY INTO <location> ✔

D. SAVE

Correct Answer: C. COPY INTO <location>

Q41. Default CSV compression on unload?

A. NONE

B. AUTO ✔

C. GZIP

D. BZIP2

Correct Answer: B. AUTO

Q42. Which object tracks file load history?

A. ACCOUNT_USAGE

B. LOAD_HISTORY

C. QUERY_HISTORY

D. COPY_HISTORY ✔

Correct Answer: D. COPY_HISTORY

Q43. What does PURGE=TRUE do?

A. Deletes source files ✔
B. Deletes target rows
C. Deletes stage
D. Deletes table
Correct Answer: A. Deletes source files

Q44. Which file type supports schema evolution?

A. CSV
B. JSON ✔
C. XML
D. Parquet
Correct Answer: B. JSON

Q45. External stages require:

A. Storage integration ✔
B. Warehouse
C. Resource monitor
D. Schema
Correct Answer: A. Storage integration

Q46. Which command lists stage files?

A. SHOW FILES
B. LIST @stage ✔
C. DESCRIBE STAGE
D. GET FILES
Correct Answer: B. LIST @stage

Q47. What is Snowpipe Streaming?

A. Batch Snowpipe
B. Low-latency ingestion ✔
C. External table refresh
D. COPY optimization
Correct Answer: B. Low-latency ingestion

Q48. Which role manages integrations?

A. SYSADMIN
B. SECURITYADMIN

C. ACCOUNTADMIN ✔

D. ORGADMIN

Correct Answer: C. ACCOUNTADMIN

Q49. What format supports nested structures?

A. CSV

B. XML

C. JSON ✔

D. TXT

Correct Answer: C. JSON

Q50. Which object stores unstructured data?

A. Table

B. External table

C. Stage ✔

D. Stream

Correct Answer: C. Stage

SECTION 4: Performance & Optimization (51–65)

Q51. What improves filter pruning?

A. Index

B. Clustering key ✔

C. Primary key

D. Constraint

Correct Answer: B. Clustering key

Q52. What does automatic clustering do?

A. Rewrites queries

B. Reorganizes micro-partitions ✔

C. Creates indexes

D. Caches results

Correct Answer: B. Reorganizes micro-partitions

Q53. Which cache stores query results?

A. Data cache

B. Metadata cache

C. Result cache ✔

D. Warehouse cache

Correct Answer: C. Result cache

Q54. When is result cache invalidated?

A. Warehouse resize

B. Table change ✔

C. Role change

D. Query timeout

Correct Answer: B. Table change

Q55. Which metric shows scanned data?

A. Execution time

B. Bytes scanned ✔

C. Rows returned

D. Credits used

Correct Answer: B. Bytes scanned

Q56. What does pruning reduce?

A. Storage size

B. Network latency

C. Micro-partitions scanned ✔

D. Credit rate

Correct Answer: C. Micro-partitions scanned

Q57. Best use of clustering keys?

A. Small tables

B. Frequently filtered columns ✔

C. Temporary tables

D. External tables

Correct Answer: B. Frequently filtered columns

Q58. Which feature prevents hot partitions?

A. Streams

B. Auto-clustering ✔

C. Time Travel

D. Tasks

Correct Answer: B. Auto-clustering

Q59. Query Profile is used for:

A. Cost estimation
B. Performance troubleshooting ✔
C. Access control
D. Encryption
Correct Answer: B. Performance troubleshooting

Q60. What shows query execution steps?

A. Query History
B. Query Profile ✔
C. Access History
D. Load History
Correct Answer: B. Query Profile

Q61. What improves join performance?

A. Broadcast join
B. Clustering
C. Warehouses ✔
D. Time Travel
Correct Answer: C. Warehouses

Q62. Which operation is metadata-only?

A. SELECT
B. CREATE TABLE ✔
C. INSERT
D. DELETE
Correct Answer: B. CREATE TABLE

Q63. Micro-partitions are:

A. User-defined
B. Fixed-size
C. Automatically managed ✔
D. Indexed
Correct Answer: C. Automatically managed

Q64. What reduces query latency most?

A. Larger warehouse
B. Result cache ✔

C. Clustering

D. Time Travel

Correct Answer: B. Result cache

Q65. What causes excessive credit usage?

A. Long-running warehouses ✔

B. Result cache

C. Metadata queries

D. Auto-suspend

Correct Answer: A. Long-running warehouses

SECTION 5: Security & Data Protection (66–85)

Q66. Which role manages users?

A. SYSADMIN

B. SECURITYADMIN ✔

C. ACCOUNTADMIN

D. PUBLIC

Correct Answer: B. SECURITYADMIN

Q67. Which feature masks sensitive data?

A. Secure views

B. Masking policies ✔

C. Network policy

D. Streams

Correct Answer: B. Masking policies

Q68. Row-level security uses:

A. Grants

B. Secure views

C. Row access policies ✔

D. Constraints

Correct Answer: C. Row access policies

Q69. Time Travel retention depends on:

A. Warehouse size

B. Table type

C. Edition ✔

D. Role

Correct Answer: C. Edition

Q70. Fail-safe duration is:

A. 1 day

B. 7 days

C. 14 days ✔

D. 30 days

Correct Answer: C. 14 days

Q71. Which object supports Time Travel?

A. Temporary tables

B. External tables

C. Permanent tables ✔

D. Streams

Correct Answer: C. Permanent tables

Q72. What enforces IP restrictions?

A. Resource Monitor

B. Network Policy ✔

C. Secure View

D. Role

Correct Answer: B. Network Policy

Q73. MFA is configured using:

A. Warehouse

B. Security Integration ✔

C. Resource Monitor

D. Task

Correct Answer: B. Security Integration

Q74. Secure views prevent:

A. Query execution

B. Data replication

C. Data leakage ✔

D. Storage usage

Correct Answer: C. Data leakage

Q75. What logs object access?

A. QUERY_HISTORY
B. ACCESS_HISTORY ✔
C. LOGIN_HISTORY
D. LOAD_HISTORY
Correct Answer: B. ACCESS_HISTORY

Q76. Which function returns current role?

A. CURRENT_USER()
B. CURRENT_ROLE() ✔
C. CURRENT_SESSION()
D. CURRENT_ACCOUNT()
Correct Answer: B. CURRENT_ROLE()

Q77. Data sharing requires:

A. Copying data
B. Same cloud provider
C. Secure share ✔
D. Replication
Correct Answer: C. Secure share

Q78. Which feature shares data read-only?

A. Streams
B. Tasks
C. Secure shares ✔
D. Stages
Correct Answer: C. Secure shares

Q79. Which role manages grants?

A. SYSADMIN
B. SECURITYADMIN ✔
C. ACCOUNTADMIN
D. PUBLIC
Correct Answer: B. SECURITYADMIN

Q80. What is VPS?

A. Virtual warehouse
B. Private Snowflake deployment ✔

C. External stage

D. Compute isolation

Correct Answer: B. Private Snowflake deployment

Q81. Who manages organizations?

A. ACCOUNTADMIN

B. SYSADMIN

C. ORGADMIN ✔

D. SECURITYADMIN

Correct Answer: C. ORGADMIN

Q82. Which feature audits user access?

A. Streams

B. Tasks

C. Access History ✔

D. Load History

Correct Answer: C. Access History

Q83. What hides column values?

A. Clustering

B. Masking policy ✔

C. Secure view

D. Encryption

Correct Answer: B. Masking policy

Q84. Which policy filters rows?

A. Masking

B. Row Access ✔

C. Network

D. Resource

Correct Answer: B. Row Access

Q85. Data is encrypted:

A. At rest only

B. In transit only

C. At rest and in transit ✔

D. On demand

Correct Answer: C. At rest and in transit

SECTION 6: Streams, Tasks & Misc (86–100)

Q86. Streams track:

A. Schema changes
B. Row-level changes ✔
C. Query changes
D. Role changes
Correct Answer: B. Row-level changes

Q87. Tasks are used for:

A. Scheduling SQL ✔
B. Streaming data
C. Sharing data
D. Query caching
Correct Answer: A. Scheduling SQL

Q88. Tasks can be triggered by:

A. Time only
B. Streams only
C. Time or stream ✔
D. User action
Correct Answer: C. Time or stream

Q89. What does MAX_RETRIES control?

A. Warehouse restarts
B. Task retry attempts ✔
C. Load retries
D. Query retries
Correct Answer: B. Task retry attempts

Q90. What suspends task execution?

A. Warehouse suspend
B. Task suspend ✔
C. Database suspend
D. Role revoke
Correct Answer: B. Task suspend

Q91. What is FLATTEN used for?

A. Compress data
B. Expand arrays ✔
C. Encrypt JSON
D. Validate data
Correct Answer: B. Expand arrays

Q92. Which function parses JSON text?

A. TO_JSON
B. CAST
C. PARSE_JSON ✔
D. FLATTEN
Correct Answer: C. PARSE_JSON

Q93. Which clause returns random rows?

A. RANDOM()
B. SAMPLE ✔
C. AVERAGE
D. LIMIT
Correct Answer: B. SAMPLE

Q94. SAMPLE with REPEATABLE ensures:

A. Faster queries
B. Same rows returned ✔
C. Larger samples
D. Sorted output
Correct Answer: B. Same rows returned

Q95. Which function detects language?

A. CORTEX.TRANSLATE
B. CORTEX.DETECT_LANGUAGE ✔
C. CORTEX.SENTIMENT
D. CORTEX.CLASSIFY
Correct Answer: B. CORTEX.DETECT_LANGUAGE

Q96. Which function summarizes text?

A. CORTEX.TRANSLATE
B. CORTEX.COMPLETE

C. CORTEX.SUMMARIZE ✔

D. CORTEX.SEARCH

Correct Answer: C. CORTEX.SUMMARIZE

Q97. Which connector supports MFA token caching?

A. JDBC ✔

B. Go

C. Python ✔

D. Node.js

Correct Answer: A. JDBC | C. Python

Q98. Which feature supports zero-copy cloning?

A. Streams

B. Tasks

C. Time Travel

D. Cloning ✔

Correct Answer: D. Cloning

Q99. Cloning copies:

A. Data only

B. Metadata only

C. Metadata + pointers ✔

D. Physical files

Correct Answer: C. Metadata + pointers

Part 2 — Snowflake Scenario-Based Practice Questions

Source files: snowflake_test.pdf (with explanations) and certification_real.pdf (same question set)

Q1. Why would the REPEATABLE keyword be used in a SAMPLE clause?

A. To limit the sample to a specific number of rows

B. To ensure no duplicates are returned when a query is run

C. To ensure the same sample is returned each time a query is run ✔

D. To force the sampling to be performed on blocks rather than rows

Correct Answer: C. To ensure the same sample is returned each time a query is run

Explanation: REPEATABLE ensures deterministic sampling using a fixed seed.

Q2. Which Snowflake feature supports the use of natural language queries to analyze unstructured data?

A. Snowsight

B. Snowflake ML

C. Snowflake Cortex ✔

D. Snowflake Horizon Catalog

Correct Answer: C. Snowflake Cortex

Explanation: Snowflake Cortex enables LLM-powered natural language processing.

Q3. Which should be used for near real-time feedback on data ingestion and ensuring the data is almost instantly available for analysis?

A. Snowpipe

B. COPY INTO <table>

C. Streams

D. Snowpipe Streaming ✔

Correct Answer: D. Snowpipe Streaming

Explanation: Snowpipe Streaming supports low-latency, near real-time ingestion.

Q4. If Role_A owns Role_B, what permissions does Role_A have based on this relationship?

A. Role_A automatically inherits all privileges assigned to Role_B

B. Role_A can use the privileges of Role_B, if Role_A is higher than Role_B in a defined hierarchy ✔

C. Role_A can grant or revoke access to the role

D. Role_A can execute any tasks assigned to Role_B without an explicit grant

Correct Answer: B. Role_A can use the privileges of Role_B, if Role_A is higher than Role_B in a defined hierarchy

Explanation: Higher roles inherit privileges from lower roles in the hierarchy.

Q5. How can a transient table be converted to a permanent table?

A. ALTER TABLE … RENAME

B. CREATE TABLE AS SELECT (CTAS)

C. ALTER TABLE … SWAP WITH ✔

D. CREATE TABLE ... CLONE

Correct Answer: C. ALTER TABLE ... SWAP WITH

Explanation: SWAP WITH exchanges metadata, effectively converting table types.

Q6. Which commands will indicate if Multi-Factor Authentication (MFA) has been enabled for a given user?

A. SHOW USERS ✔

B. DESCRIBE USER ✔

C. SHOW PARAMETERS

D. SHOW USER FUNCTIONS

E. SHOW CONNECTIONS

Correct Answer: A. SHOW USERS | B. DESCRIBE USER

Explanation: MFA status is visible via user metadata.

Q7. What will happen when the data retention period for an object is set to 0 days?

A. Streams containing the object will become stale

B. Time Travel will be deactivated for the object ✔

C. Data replication will be disabled for the object

D. Data sharing will be disabled for the object

Correct Answer: B. Time Travel will be deactivated for the object

Explanation: Zero retention disables Time Travel.

Q8. What type of queries will benefit from the Search Optimization Service when using equality predicates?

A. Queries using point-lookups ✔

B. Queries using data substrings

C. Queries with geospatial data

D. Queries of semi-structured data

Correct Answer: A. Queries using point-lookups

Explanation: Search Optimization accelerates point-lookups using = filters.

Q9. A user has an issue with queries taking too long to start execution during peak hours. Which step should the user take?

A. Reduce the size of the virtual warehouse

B. Increase the size of the virtual warehouse

C. Disable auto-scale mode for the warehouse

D. Convert the warehouse to a multi-cluster warehouse ✔

Correct Answer: D. Convert the warehouse to a multi-cluster warehouse

Explanation: Multi-cluster warehouses reduce queuing during concurrency spikes.

Q10. A user needs to know: 1. Who accessed the data 2. When the data was accessed 3. What columns were accessed Which ACCOUNT_USAGE schema view should be used?

A. QUERY_HISTORY

B. ACCESS_HISTORY ✔

C. EVENT_USAGE_HISTORY

D. OBJECT_DEPENDENCIES

Correct Answer: B. ACCESS_HISTORY

Q11. Which command or parameter is used to execute a task every 30 seconds?

A. WHEN

B. SCHEDULE ✔

C. USER_TASK_TIMEOUT_MS

D. EXECUTE TASK

Correct Answer: B. SCHEDULE

Explanation: Tasks are scheduled using the SCHEDULE parameter (cron or interval).

Q12. Which queries need an active warehouse to run? (Select TWO)

A. SELECT COUNT(*) FROM <TABLE_NAME>;

B. CREATE TABLE <TABLE_NAME> (<COLUMN> <DATATYPE>);

C. DESCRIBE TABLE <TABLE_NAME>;

D. DELETE FROM <TABLE_NAME>; ✔

E. SELECT * FROM <TABLE_NAME>; ✔

Correct Answer: D. DELETE FROM <TABLE_NAME>; | E. SELECT * FROM <TABLE_NAME>;

Explanation: DML and SELECT queries require compute; metadata operations do

Q13. Bulk data is loaded daily from an external stage using the following command: COPY INTO source_table FROM s3://source/data/files STORAGE_INTEGRATION = source_int FILE_FORMAT = (FORMAT_NAME = source_format) PATTERN = 'sales.*\\.csv'; How will the data be loaded?

A. All files that match the format will be appended

B. Only new files will be loaded; previously loaded files will be skipped ✔

C. All CSV files in the path will be loaded

D. The table will be truncated before loading

Correct Answer: B. Only new files will be loaded; previously loaded files will be skipped

Explanation: Snowflake tracks file load history and skips previously loaded files.

Q14. Which Snowflake service is managed within the Cloud Services layer?

A. Data storage

B. File compression

C. Query processing

D. Infrastructure management ✔

Correct Answer: D. Infrastructure management

Explanation: Cloud Services handles metadata, security, and infrastructure

Q15. Which data sharing mechanism can be used to share data privately or publicly using the Snowflake Marketplace?

A. Listing ✔

B. Direct share

C. Reader account

D. Data Exchange

Correct Answer: A. Listing

Explanation: Listings are used to publish datasets on the Marketplace.

Q16. What does the : operator do when used with a VARIANT value?

A. Converts a VARIANT into an object

B. Extracts the raw value from a VARIANT column ✔

C. Casts the value to a specified data type

D. Generates a random sample

Correct Answer: B. Extracts the raw value from a VARIANT column

Explanation: The colon operator extracts elements from VARIANT data.

Q17. Which feature is used to protect sensitive data at query run time?

A. Object tagging

B. Data encryption

C. User-Defined Functions (UDFs)

D. Dynamic Data Masking ✔

Correct Answer: D. Dynamic Data Masking

Explanation: Dynamic Data Masking hides data dynamically based on role.

Q18. When a replication schedule is set, how does Snowflake manage refreshes for a secondary object?

A. Multiple refreshes will run concurrently

B. Only one refresh will run at a time ✔

C. Refreshes run regardless of role privileges

D. Delayed refresh cancels the next scheduled refresh

Correct Answer: B. Only one refresh will run at a time

Explanation: Snowflake ensures only one replication refresh runs at a time.

Q19. A transient table is created with clustering, masking, and row access policy. Which feature will incur both compute and storage costs?

A. Row access policy

B. Automatic Clustering ✔

C. Dynamic Data Masking policy

D. NOT NULL constraint

Correct Answer: B. Automatic Clustering

Q20. A table is created with clustering, a masking policy, and a row access policy. Which feature of this table will incur both compute and storage costs?

A. The row access policy

B. The Automatic Clustering ✔

C. The Dynamic Data Masking policy

D. The NOT NULL constraint

Correct Answer: B. The Automatic Clustering

Q21. When a replication schedule is set, how does Snowflake manage refreshes for failover groups?

A. Multiple refreshes will run concurrently

B. Only one refresh will run at a time ✔

C. Refreshes will run regardless of role privileges

D. If the current refresh is delayed, the next scheduled refresh will be canceled

Correct Answer: B. Only one refresh will run at a time

Q22. Which command automates the ingestion of data when it arrives in a stage?

A. PUT

B. COPY FILES

C. CREATE PIPE ✔

D. COPY INTO <table>

Correct Answer: C. CREATE PIPE

Q23. When resizing a currently running virtual warehouse, what happens to in-progress queries?

A. In-progress queries will be canceled

B. Queries will continue running using the new warehouse size

C. Queries will continue running using the current warehouse size ✔

D. Queries will be paused and reprocessed using the new size

Correct Answer: C. Queries will continue running using the current warehouse size

Q24. Which table type is optimized for transactional workloads?

A. Event

B. Hybrid ✔

C. Dynamic

D. Permanent

Correct Answer: B. Hybrid

Q25. Which data sharing mechanism can be used to share data privately or publicly within the Snowflake Marketplace?

A. Listing ✔

B. Direct share

C. Reader account

D. Data Exchange

Correct Answer: A. Listing

Q26. Which function, when combined with an INSERT statement, will insert semi-structured data into a VARIANT column?

A. TO_JSON

B. TO_VARCHAR

C. PARSE_JSON ✔

D. OBJECT_INSERT

Correct Answer: C. PARSE_JSON

Q27. What is the MINIMUM Snowflake edition required to create a materialized view?

A. Standard

B. Enterprise ✔

C. Business Critical

D. Virtual Private Snowflake (VPS)

Correct Answer: B. Enterprise

Q28. Which Snowflake feature or service can be used with a standard materialized view?

A. Time Travel ✔

B. Fail-safe

C. Clustering

D. Data sharing

Correct Answer: A. Time Travel

Q29. Which action allows users to directly define the physical organization of data to optimize query performance on large tables?

A. Manually configuring micro-partitions

B. Specifying clustering keys on columns frequently used in filters ✔

C. Using PARTITION BY clauses

D. Writing scripts to reorder data files

Correct Answer: B. Specifying clustering keys on columns frequently used in filters

Q30. When working with a secure view that will be shared with other Snowflake accounts, which functions return a NULL value? (Select TWO)

A. CURRENT_ACCOUNT_NAME() ✔

B. CURRENT_ROLE()

C. CURRENT_ROLE_TYPE()

D. CURRENT_SESSION() ✔

E. CURRENT_USER()

Correct Answer: A. CURRENT_ACCOUNT_NAME() | D. CURRENT_SESSION()

Q31. A query runs significantly slower than expected. Which Query Profile metrics help identify the issue? (Select TWO)

A. Processing

B. Resource usage

C. Query Plan execution steps ✔

D. Bytes scanned

E. Partitions scanned ✔

Correct Answer: C. Query Plan execution steps | E. Partitions scanned

Q32. What is the FLATTEN function used for when working with semi-structured data in Snowflake?

A. Extract and expand nested arrays or objects into relational format ✔

B. Compress semi-structured data

C. Convert structured data into semi-structured

D. Aggregate semi-structured columns

Correct Answer: A. Extract and expand nested arrays or objects into relational format

Q33. Where can file formats be specified when unloading data from Snowflake? (Select TWO)

A. SELECT statement

B. CREATE STAGE statement ✔

C. CREATE DATABASE statement

D. CREATE INTEGRATION statement

E. COPY INTO <location> ✔

Correct Answer: B. CREATE STAGE statement | E. COPY INTO <location>

Q34. When unloading data from a query into a location, which VALIDATION_MODE parameter returns the query results?

A. RETURN_<n>_ROWS

B. RETURN_ROWS ✔

C. RETURN_ERRORS

D. RETURN_ALL_ERRORS

Correct Answer: B. RETURN_ROWS

Q35. What is the default compression option when unloading CSV data using COPY INTO <location>?

A. BZ2

B. AUTO ✔

C. NONE

D. DEFLATE

Correct Answer: B. AUTO

Q36. A user wants to temporarily define access to a staged file without the need for the user to have stage-level privileges. Which type of Snowflake URL should be used?

A. Pre-signed

B. Scoped ✔

C. External stage

D. File

Correct Answer: B. Scoped

Q37. Taking which step will MINIMIZE the number of queries waiting to run on a virtual warehouse?

A. Increase the WAREHOUSE_SIZE

B. Increase the value of the AUTO_SUSPEND parameter

C. Set the WAIT_FOR_COMPLETION parameter to FALSE

D. Increase the value of the MAX_CLUSTER_COUNT ✔

Correct Answer: D. Increase the value of the MAX_CLUSTER_COUNT

Q38. Using what type of table will help users MINIMIZE storage costs for tables that contain non-critical data and change frequently?

A. Temporary table

B. Transient table ✔

C. Dynamic table

D. Hybrid table

Correct Answer: B. Transient table

Q39. Which feature helps improve query performance for workloads that use common predicate patterns?

A. Auto-suspend

B. Materialized view

C. Query acceleration service

D. Search optimization service ✔

Correct Answer: D. Search optimization service

Q40. How can a relational table be unloaded into a JSON file?

A. Use the OBJECT_CONSTRUCT function with the COPY INTO <location> command

B. Use the COPY INTO <location> command with the file_format set as JSON ✔

C. Use the PUT command with the file_format set as JSON

D. Use the GET command with the file_format set as JSON

Correct Answer: B. Use the COPY INTO <location> command with the file_format set as JSON

Q41. A virtual warehouse has the same value for MIN_CLUSTER_COUNT and MAX_CLUSTER_COUNT. What will happen when the warehouse is started?

A. Snowflake will start all available clusters

B. Snowflake will start enough clusters to handle the current workload

C. Snowflake will start a single cluster and start additional clusters after 6 minutes

D. Snowflake will start the fixed number of clusters defined by MIN_CLUSTER_COUNT ✔

Correct Answer: D. Snowflake will start the fixed number of clusters defined by MIN_CLUSTER_COUNT

Q42. By default, when is a resource monitor reset to 0?

A. At the beginning of each day

B. At the beginning of each week

C. At the beginning of each month ✔

D. At the beginning of each year

Correct Answer: C. At the beginning of each month

Q43. Which uses of directory tables are valid? (Select TWO)

A. Retrieve a list of all unstructured files on a stage ✔

B. Sync and automatically refresh metadata for directory tables on external and internal stages ✔

C. Remove files from a directory table using the GET command

D. Create a view for unstructured data

E. Process unstructured data using Snowflake system-defined functions

Correct Answer: A. Retrieve a list of all unstructured files on a stage | B. Sync and automatically refresh metadata for directory tables on external and internal stages

Q44. How should a Snowflake user access a third-party SaaS service to process unstructured data?

A. Use internal functions

B. Use external functions ✔

C. Use process functions

D. Use an API gateway

Correct Answer: B. Use external functions

Q45. What can a directory table be used for? (Select TWO)

A. Retrieve a list of all unstructured files on a stage ✔

B. Set an automatic refresh of the metadata for a directory table on internal and external stages ✔

C. Remove files from a directory table using the GET command

D. Create a view for unstructured data

E. Process unstructured data using Snowflake system-defined functions

Correct Answer: A. Retrieve a list of all unstructured files on a stage | B. Set an automatic refresh of the metadata for a directory table on internal and external stages

Q46. Which feature allows a user to review the monthly projected costs of a Snowflake resource against the monthly budget?

A. Budgets ✔

B. Auto-suspend

C. Scaling Policy

D. Resource monitor

Correct Answer: A. Budgets

Q47. Which setting will help control credit consumption for a virtual warehouse that needs to run in auto-scale mode?

A. Set the warehouse scaling policy to Economy ✔

B. Set the MIN_CLUSTER_COUNT to 1

C. Set the AUTO_SUSPEND property of the warehouse to 0

D. Set the MIN_CLUSTER_COUNT to 2 and set the MAX_CLUSTER_COUNT to 2

Correct Answer: A. Set the warehouse scaling policy to Economy

Q48. When unloading data, which file format will keep floating-point number column data without truncating it?

A. XML

B. JSON

C. Parquet ✔

D. CSV

Correct Answer: C. Parquet

Q49. What is the primary purpose for configuring multiple Identity Providers (IdPs) in Snowflake?

A. To enforce complex password rules

B. To improve the performance of a query

C. To manage encryption keys for secure data storage

D. To allow different authentication methods for multiple user groups ✔

Correct Answer: D. To allow different authentication methods for multiple user groups

Q50. A task is set to MAX_RETRIES = 3 and a virtual warehouse is set to AUTO_SUSPEND = 60. How will the task perform with this configuration?

A. The task will attempt to run three times, and the warehouse will not automatically suspend

B. The task will attempt to run indefinitely, and the warehouse will suspend after 60 seconds of inactivity

C. The task will run once, and the warehouse will suspend immediately after the task has finished

D. The task will attempt to run three times, and the warehouse will suspend after 60 seconds of inactivity ✔

Correct Answer: D. The task will attempt to run three times, and the warehouse will suspend after 60 seconds of inactivity

Q51. What is the recommended way to disable Multi-Factor Authentication (MFA) for a Snowflake user?

A. Turn off the ALLOW_CLIENT_MFA_CACHING parameter

B. Set the MINS_TO_BYPASS_MFA = 30 parameter at the user level

C. Set DISABLE_MFA = TRUE at the user level ✔

D. Temporarily pause the approved MFA methods

Correct Answer: C. Set DISABLE_MFA = TRUE at the user level

Q52. What should a Snowflake user do when assigning system-defined roles?

A. Assign the SYSADMIN role only to project managers

B. Assign the ACCOUNTADMIN role to a minimum of two users ✔

C. Assign the SECURITYADMIN role to only one user per organization

D. Grant the ORGADMIN role to at least one user for each account in each region

Correct Answer: B. Assign the ACCOUNTADMIN role to a minimum of two users

Q53. What is the effect of using FORCE = TRUE in a COPY INTO <table> command?

A. All files will be loaded, even ones with a load status of UNKNOWN

B. Files that have already been successfully loaded into the table will be loaded again ✔

C. Files that have already been successfully loaded into the table will not be loaded again

D. String columns longer than the defined column length will be loaded anyway

Correct Answer: B. Files that have already been successfully loaded into the table will be loaded again

Q54. What is the effect of using FORCE = TRUE in a COPY INTO <table> command?

A. All files will be loaded, even ones with a load status of UNKNOWN

B. Files that have already been successfully loaded into the table will be loaded again ✔

C. Files that have already been successfully loaded into the table will not be loaded again

D. String columns longer than the column definition length will be loaded anyway

Correct Answer: B. Files that have already been successfully loaded into the table will be loaded again

Q55. Which statement will return all rows from table_A?

A. SELECT * FROM table_A TABLESAMPLE (0);

B. SELECT * FROM table_A TABLESAMPLE (100); ✔

C. SELECT * FROM table_A SAMPLE (10);

D. SELECT * FROM table_A TABLESAMPLE BERNOULLI (10);

Correct Answer: B. SELECT * FROM table_A TABLESAMPLE (100);

Q56. Which function converts a VARIANT 'null' value to a SQL NULL?

A. CAST

B. TO_CHAR

C. TRY_CAST ✔

D. TO_BINARY

Correct Answer: C. TRY_CAST

Q57. Which function, when added to a SELECT statement, will return a randomly selected, specified number of rows from a table?

A. AVERAGE(<num> ROWS)

B. MEDIAN(<num> ROWS)

C. SAMPLE(<num> ROWS) ✔

D. RANDOM(<num> ROWS)

Correct Answer: C. SAMPLE(<num> ROWS)

Q58. Which Snowflake edition supports a dedicated metadata store and separate compute clusters that are inaccessible to other Snowflake customers?

A. Standard

B. Enterprise

C. Business Critical

D. Virtual Private Snowflake (VPS) ✔

Correct Answer: D. Virtual Private Snowflake (VPS)

Q59. How can data be shared in Snowflake with another account located in the same region without copying or moving the data?

A. Use the Snowflake Marketplace

B. Use a stream

C. Use a Data Exchange

D. Use a direct share ✔

Correct Answer: D. Use a direct share

Q60. Which connectors support Multi-Factor Authentication (MFA) token caching when connecting to Snowflake? (Select TWO)

A. JDBC ✔

B. Go

C. PHP

D. Node.js

E. Python ✔

Correct Answer: A. JDBC | E. Python

Part 3 — Snowflake Concept & Exam-Angle Flashcard MCQs

Source file: Snowflake Certification.pdf (concept slides with embedded MCQs; duplicated slides removed)

Q1. What key limitation of traditional DBs does Snowflake solve?

A. SQL support
B. Tight coupling of compute & storage ✔
C. Authentication
D. Data types
Correct Answer: B. Tight coupling of compute & storage

Q2. Which layer performs query optimization?

A. Storage
B. Warehouse
C. Cloud Services ✔
D. Client
Correct Answer: C. Cloud Services

Q3. Cloud Services does NOT handle:

A. Metadata
B. Query parsing
C. Query execution ✔
D. Security
Correct Answer: C. Query execution

Q4. Which component directly consumes credits?

A. Cloud Services
B. Storage
C. Virtual Warehouse ✔
D. Metadata
Correct Answer: C. Virtual Warehouse

Q5. How many copies of data exist for multiple warehouses?

A. One per warehouse
B. One per user
C. Single shared copy ✔
D. Depends on size
Correct Answer: C. Single shared copy

Q6. Which statement is TRUE?

A. Warehouses store data

B. Warehouses execute queries ✔

C. Warehouses manage metadata

D. Warehouses control access

Correct Answer: B. Warehouses execute queries

Q7. How does Snowflake handle concurrent workloads?

A. Locks tables

B. Queues queries

C. Uses separate warehouses ✔

D. Replicates data

Correct Answer: C. Uses separate warehouses

Q8. Which feature depends on edition?

A. Warehouses

B. SQL

C. Time Travel retention ✔

D. Result cache

Correct Answer: C. Time Travel retention

Q9. What architectural feature enables elasticity?

A. Caching

B. Clustering

C. Decoupled compute & storage ✔

D. Replication

Correct Answer: C. Decoupled compute & storage

Q10. What does a virtual warehouse primarily provide?

A. Metadata

B. Storage

C. Compute resources ✔

D. Access control

Correct Answer: C. Compute resources

Q11. Which is NOT true about a warehouse?

A. Stateless

B. Independent

C. Stores table data ✔

D. Consumes credits

Correct Answer: C. Stores table data

Q12. Scaling a warehouse from M to L mainly increases:

A. Storage

B. Number of users

C. Compute power ✔

D. Metadata

Correct Answer: C. Compute power

Q13. Dashboard with 300 concurrent users → best solution?

A. XL warehouse

B. Increase storage

C. Multi-cluster warehouse ✔

D. Result cache only

Correct Answer: C. Multi-cluster warehouse

Q14. Each cluster in a multi-cluster warehouse:

A. Shares cache

B. Shares credits

C. Consumes credits independently ✔

D. Stores data

Correct Answer: C. Consumes credits independently

Q15. Which feature immediately reduces wasted credits?

A. Result cache

B. Auto-suspend ✔

C. Clustering

D. Replication

Correct Answer: B. Auto-suspend

Q16. Warehouse is suspended. Same query is re-run. What happens?

A. Query fails

B. Full re-scan

C. Uses result cache ✔

D. Needs manual resume

Correct Answer: C. Uses result cache

Q17. Which cache is cleared when a warehouse suspends?

A. Result cache
B. Metadata cache
C. Warehouse cache ✔
D. Cloud Services cache
Correct Answer: C. Warehouse cache

Q18. Best practice for ETL + BI workloads?

A. One large warehouse
B. Multi-cluster only
C. Separate warehouses ✔
D. Disable caching
Correct Answer: C. Separate warehouses

Q19. Which factor directly increases credit usage?

A. Table size
B. Query complexity
C. Warehouse size & runtime ✔
D. Result cache
Correct Answer: C. Warehouse size & runtime

Q20. Which statement is ALWAYS true?

A. Warehouses store data
B. Warehouses share cache
C. Warehouses consume credits only when running ✔
D. Warehouses manage security
Correct Answer: C. Warehouses consume credits only when running

Q21. Which statement about Snowflake storage is TRUE?

A. Each warehouse has its own storage
B. Storage is shared across warehouses ✔
C. Users define partitions
D. Storage consumes compute credits
Correct Answer: B. Storage is shared across warehouses

Q22. Who manages Snowflake storage?

A. DBA
B. User
C. Snowflake automatically ✔
D. Cloud provider only
Correct Answer: C. Snowflake automatically

Q23. Micro-partitions are:

A. User-defined
B. Fixed-size row blocks
C. Automatically managed units of storage ✔
D. Warehouse-specific
Correct Answer: C. Automatically managed units of storage

Q24. Columnar storage improves performance mainly by:

A. Increasing warehouse size
B. Scanning only required columns ✔
C. Adding indexes
D. Reordering rows
Correct Answer: B. Scanning only required columns

Q25. What enables Snowflake to skip scanning data?

A. Indexes
B. Warehouse cache
C. Micro-partition metadata ✔
D. Result cache
Correct Answer: C. Micro-partition metadata

Q26. Which scenario benefits most from pruning?

A. Small dimension table
B. Large fact table with date filters ✔
C. Temporary tables
D. Lookup tables
Correct Answer: B. Large fact table with date filters

Q27. Partition pruning primarily reduces:

A. Storage size
B. Number of users
C. Data scanned ✔

D. Metadata size

Correct Answer: C. Data scanned

Q28. Clustering improves performance by:

A. Creating indexes

B. Improving partition pruning ✔

C. Increasing cache size

D. Sorting rows permanently

Correct Answer: B. Improving partition pruning

Q29. Manual clustering requires:

A. No compute

B. Snowflake Support

C. Compute credits for reclustering ✔

D. Enterprise edition only

Correct Answer: C. Compute credits for reclustering

Q30. Automatic clustering:

A. Is free

B. Requires user scheduling

C. Is managed by Snowflake ✔

D. Replaces pruning

Correct Answer: C. Is managed by Snowflake

Q31. Clustering is NOT recommended for:

A. Large fact tables

B. Frequently filtered columns

C. Small lookup tables ✔

D. Date-based filters

Correct Answer: C. Small lookup tables

Q32. Which feature does Snowflake NOT support?

A. Partition pruning

B. Clustering

C. Indexes ✔

D. Metadata

Correct Answer: C. Indexes

Q33. Best way to improve performance on large filtered table?

A. Add index
B. Increase storage
C. Improve pruning via clustering ✔
D. Disable caching
Correct Answer: C. Improve pruning via clustering

Q34. Which layer parses and optimizes SQL?

A. Storage
B. Warehouse
C. Cloud Services ✔
D. Client tool
Correct Answer: C. Cloud Services

Q35. Which condition INVALIDATES result cache?

A. Same query text
B. Same role
C. Data change in table ✔
D. Warehouse suspension
Correct Answer: C. Data change in table

Q36. Which is evaluated first when a query is submitted?

A. Warehouse cache
B. Storage
C. Result cache ✔
D. Metadata cache
Correct Answer: C. Result cache

Q37. A warehouse is suspended. Same SELECT query runs again. What happens?

A. Query fails
B. Warehouse resumes
C. Result cache is used ✔
D. Full data scan
Correct Answer: C. Result cache is used

Q38. Metadata cache is primarily used for:

A. Query execution
B. Result storage

C. Optimization & pruning ✔

D. Cost tracking

Correct Answer: C. Optimization & pruning

Q39. Which cache survives warehouse suspension?

A. Warehouse cache

B. Result cache ✔

C. Temporary tables

D. Local disk cache

Correct Answer: B. Result cache

Q40. Which feature allows queries to return results without using compute?

A. Clustering

B. Result cache ✔

C. Warehouse cache

D. Time Travel

Correct Answer: B. Result cache

Q41. Which two methods are used to load data into Snowflake?

A. INSERT and UPDATE

B. COPY INTO and Snowpipe ✔

C. PUT and GET

D. Streams and Tasks

Correct Answer: B. COPY INTO and Snowpipe

Q42. Why are stages required in Snowflake?

A. To store metadata

B. To temporarily hold data files before loading ✔

C. To cache query results

D. To manage access control

Correct Answer: B. To temporarily hold data files before loading

Q43. Which is TRUE about internal stages?

A. Require AWS/Azure credentials

B. Managed by Snowflake ✔

C. Store data permanently in tables

D. Only support Snowpipe

Correct Answer: B. Managed by Snowflake

Q44. External stages are primarily used to:

A. Store data in Snowflake
B. Reference cloud storage outside Snowflake ✔
C. Cache query results
D. Replace tables
Correct Answer: B. Reference cloud storage outside Snowflake

Q45. What is the default behavior of COPY INTO?

A. Reloads files every time
B. Loads each file once only ✔
C. Updates existing rows
D. Deletes source files
Correct Answer: B. Loads each file once only

Q46. Why does Snowflake track loaded files?

A. To improve compression
B. To prevent duplicate loading ✔
C. To control access
D. To manage caching
Correct Answer: B. To prevent duplicate loading

Q47. What is the purpose of a file format?

A. Store data
B. Define how files are interpreted ✔
C. Improve query performance
D. Encrypt data
Correct Answer: B. Define how files are interpreted

Q48. Snowpipe loads data using:

A. INSERT statements
B. Streams
C. COPY INTO internally ✔
D. Tasks
Correct Answer: C. COPY INTO internally

Q49. Which statement about Snowpipe is TRUE?

A. Requires running warehouse
B. Free of cost
C. Serverless and auto-scales ✔
D. Replaces COPY INTO syntax
Correct Answer: C. Serverless and auto-scales

Q50. Which loading method is best for near-real-time ingestion?

A. INSERT
B. COPY INTO
C. Snowpipe ✔
D. Clustering
Correct Answer: C. Snowpipe

Q51. Which formats are natively supported as semi-structured data?

A. CSV only
B. JSON, AVRO, PARQUET, XML ✔
C. TXT only
D. Relational tables only
Correct Answer: B. JSON, AVRO, PARQUET, XML

Q52. Which data type is used to store JSON in Snowflake?

A. STRING
B. OBJECT
C. ARRAY
D. VARIANT ✔
Correct Answer: D. VARIANT

Q53. Snowflake uses which approach for semi-structured data?

A. Schema-on-write
B. Schema-on-read ✔
C. Schema-on-update
D. No schema ever
Correct Answer: B. Schema-on-read

Q54. How are JSON keys treated in Snowflake?

A. Case-insensitive
B. Auto-converted to uppercase
C. Case-sensitive ✔

D. Converted to columns automatically

Correct Answer: C. Case-sensitive

Q55. Which structure typically requires FLATTEN?

A. Scalar value

B. Object

C. Array ✔

D. STRING

Correct Answer: C. Array

Q56. What is the primary purpose of FLATTEN?

A. Convert JSON to text

B. Compress arrays

C. Expand arrays into rows ✔

D. Enforce schema

Correct Answer: C. Expand arrays into rows

Q57. When does FLATTEN operate?

A. During data load

B. During storage

C. During query execution ✔

D. During clustering

Correct Answer: C. During query execution

Q58. Which combination is correct for querying JSON arrays?

A. STRING + JOIN

B. VARIANT + FLATTEN ✔

C. OBJECT + INDEX

D. ARRAY + CLUSTERING

Correct Answer: B. VARIANT + FLATTEN

Q59. Which features are used for data recovery in Snowflake?

A. Indexes and backups

B. Time Travel and Fail-safe ✔

C. Clustering and caching

D. Replication only

Correct Answer: B. Time Travel and Fail-safe

Q60. What can Time Travel be used for?

A. Query historical data only
B. Restore dropped objects only
C. Query and restore data ✔
D. Backup to cloud storage
Correct Answer: C. Query and restore data

Q61. Maximum Time Travel retention in Enterprise edition is:

A. 1 day
B. 7 days
C. 30 days
D. 90 days ✔
Correct Answer: D. 90 days

Q62. Which command restores a dropped table?

A. RESTORE
B. RECOVER
C. UNDROP ✔
D. CLONE
Correct Answer: C. UNDROP

Q63. Who can access data in Fail-safe?

A. ACCOUNTADMIN
B. SECURITYADMIN
C. SYSADMIN
D. Snowflake Support only ✔
Correct Answer: D. Snowflake Support only

Q64. Which feature allows users to query historical data?

A. Fail-safe
B. Replication
C. Time Travel ✔
D. Cloning
Correct Answer: C. Time Travel

Q65. What is TRUE about zero-copy cloning?

A. Data copied immediately
B. Requires extra storage upfront

C. Storage is shared initially ✔

D. Clone is read-only

Correct Answer: C. Storage is shared initially

Q66. Which feature is BEST for creating a dev copy instantly?

A. Backup

B. Replication

C. Zero-copy cloning ✔

D. Fail-safe

Correct Answer: C. Zero-copy cloning

Q67. In Snowflake, privileges are granted to:

A. Users

B. Warehouses

C. Roles ✔

D. Databases

Correct Answer: C. Roles

Q68. Which statement about roles is TRUE?

A. Roles cannot inherit

B. Users have only one role

C. Roles can be hierarchical ✔

D. PUBLIC has admin privileges

Correct Answer: C. Roles can be hierarchical

Q69. Which role is recommended for daily object creation?

A. ACCOUNTADMIN

B. SYSADMIN ✔

C. SECURITYADMIN

D. USERADMIN

Correct Answer: B. SYSADMIN

Q70. What does WITH GRANT OPTION allow?

A. Transfer ownership

B. Revoke privileges

C. Grant same privilege to others ✔

D. Enable Time Travel

Correct Answer: C. Grant same privilege to others

Q71. Which capability is unique to an object owner?

A. SELECT
B. INSERT
C. DROP ✔
D. VIEW
Correct Answer: C. DROP

Q72. When is masking policy applied?

A. During data load
B. During storage
C. During query execution ✔
D. During cloning
Correct Answer: C. During query execution

Q73. Row access policies control:

A. Columns returned
B. Rows returned ✔
C. Storage format
D. Data loading
Correct Answer: B. Rows returned

Q74. Which feature hides salary values but shows all rows?

A. Row access
B. Network policy
C. Masking policy ✔
D. Clustering
Correct Answer: C. Masking policy

Q75. Network policies are used to:

A. Mask columns
B. Filter rows
C. Restrict IP access ✔
D. Encrypt data
Correct Answer: C. Restrict IP access

Q76. Which combination best enforces data security?

A. RBAC only
B. Masking only
C. RBAC + masking + row access ✔
D. Clustering
Correct Answer: C. RBAC + masking + row access

Q77. Which role primarily manages object privileges?

A. USERADMIN
B. SYSADMIN
C. SECURITYADMIN ✔
D. PUBLIC
Correct Answer: C. SECURITYADMIN

Q78. What most directly improves Snowflake query performance?

A. Adding indexes
B. Increasing storage
C. Reducing data scanned ✔
D. Increasing Time Travel retention
Correct Answer: C. Reducing data scanned

Q79. What is the primary purpose of Query Profile?

A. Security auditing
B. Cost billing
C. Performance troubleshooting ✔
D. Backup analysis
Correct Answer: C. Performance troubleshooting

Q80. Which strategy often minimizes cost for heavy queries?

A. Small warehouse running long
B. Large warehouse running briefly ✔
C. Disable auto-suspend
D. Use result cache only
Correct Answer: B. Large warehouse running briefly

Q81. Which component usually contributes MOST to Snowflake cost?

A. Storage
B. Data transfer
C. Compute credits ✔

D. Metadata

Correct Answer: C. Compute credits

Q82. Which feature MOST effectively reduces wasted compute credits?

A. Result cache

B. Auto-suspend warehouses ✔

C. Clustering

D. Replication

Correct Answer: B. Auto-suspend warehouses

Q83. What action can a resource monitor take?

A. Delete data

B. Optimize queries

C. Suspend warehouses ✔

D. Increase storage

Correct Answer: C. Suspend warehouses

Q84. Credit usage can be analyzed at which level?

A. Table

B. Column

C. Warehouse ✔

D. File

Correct Answer: C. Warehouse

Q85. Which feature helps control runaway compute costs?

A. Clustering

B. Result cache

C. Resource monitor ✔

D. Time Travel

Correct Answer: C. Resource monitor

Q86. Which statement about Snowflake data sharing is TRUE?

A. Data is copied to the consumer

B. ETL is required

C. Data is shared without duplication ✔

D. Consumer uses provider warehouse

Correct Answer: C. Data is shared without duplication

Q87. Who pays for compute when querying shared data?

A. Provider

B. Snowflake

C. Consumer ✔

Correct Answer: C. Consumer

Q88. Who pays for storage in a reader account?

A. Consumer

B. Provider ✔

C. Snowflake

D. Shared

Correct Answer: B. Provider

Q89. Replication is primarily used for:

A. Cost optimization

B. Query performance

C. Disaster recovery ✔

D. Data loading

Correct Answer: C. Disaster recovery